

OpenAI Strategic BDR APAC Interview Prep

This is the clean prep version for live use or screen sharing. It keeps the useful answers, removes private notes, and marks facts that should be verified before use.

Use Rules

- Use this as the visible prep doc.
- Keep raw notes, internal wiki pages, private repositories, and AI tool windows closed.
- Say exact metrics only when personally confident.
- Verify current OpenAI product, model, security, privacy, and account facts before using detailed claims.
- If unsure, say: "I want to be precise rather than inflate that. The directional proof is..."

Core Positioning

One-Liner

I am an APAC enterprise GTM operator who has sold into regulated and technical buyers, then built AI-assisted GTM workflows. OpenAI is compelling because this role sits where those threads meet: turning AI interest, product usage, pilots, and developer-led experimentation into high-conviction enterprise opportunities.

Proof Pillars

Pillar	Proof
Regulated APAC GTM	Airwallex and Boxo fintech/platform selling in Singapore and APAC.
Stakeholder mapping	Business, product, engineering, legal, compliance, and executive stakeholders.
Quota discipline	Airwallex: 200%+ annual quota attainment and 22 new logos.
Enterprise/platform partnerships	Boxo: USD 450K ARR across fintech clients and 9 strategic super-app partnerships.
AI-native GTM	Salescraft/GTM Workspace: account intelligence, enrichment, scoring, outbound, CRM hygiene, and seller handoffs.

Review-first AI	Built AI workflows around sources, permissions, review gates, and measurable outputs.
-----------------	---

Tell Me About Yourself

15 Seconds

I am an APAC enterprise GTM operator who has sold into regulated and technical buyers, then built AI-assisted GTM workflows. OpenAI is compelling because this role combines enterprise account strategy, product signal, stakeholder mapping, and AI adoption.

30 Seconds

My background is enterprise GTM and AI systems. At Airwallex, I built sales discipline in fintech and exceeded target. At Boxo, I worked on platform partnerships across executive, product, engineering, legal, and compliance stakeholders. Through Salescraft and GTM Workspace, I built source-backed account-intelligence and prospecting workflows. That is why this OpenAI role feels specific: it is about turning AI interest, product signal, and developer-led experimentation into high-conviction enterprise opportunities in APAC.

2 Minutes

I have spent most of my career at the intersection of enterprise GTM, technical products, and ambiguous early markets.

At Boxo, I learned enterprise sales in an environment where the company did not yet have all the normal enterprise proof. That forced me to make deals underwriteable: map stakeholders, understand product and compliance concerns, coordinate internally with product and engineering, and keep executive conversations tied to a real implementation path. The lesson was that complex sales is structured risk removal, not just persuasion.

At Airwallex, I got the harder sales cadence: fintech buyers, quota pressure, market expansion, SDR/AE handoffs, and team management. I exceeded target and learned how much GTM performance depends on the operating system around the rep: ICP discipline, account research, qualification, follow-up, CRM hygiene, and coaching.

That led naturally into Salescraft and GTM Workspace. I started building source-backed AI workflows around the same GTM loops I had lived manually: account intelligence, enrichment, scoring, outbound quality, seller handoffs, and review gates. The broader pattern in my work has become clear: I like taking ambiguous workflows, mapping the real sources and stakeholders, and turning them into repeatable systems that humans can trust.

OpenAI's Strategic BDR role is compelling because it combines those threads. It is not just top-of-funnel volume. The job is to help turn product usage, pilots, inbound demand, and developer-led experimentation into high-conviction enterprise opportunities. That means understanding the customer, the workflow, the stakeholder map, the governance concerns, and the Account Director's path to a real deal.

Why OpenAI

OpenAI is one of the few companies where the enterprise GTM problem is as interesting as the product. There is already enormous awareness, experimentation, and developer-led usage. The hard part is helping serious organizations move from curiosity and pilots into governed deployment with a real business case. That translation layer is what attracts me: understanding the workflow, mapping the account, handling regulated-buyer concerns, and helping shape durable enterprise opportunities.

What this should prove:

- Not attracted only by brand.
- Understands enterprise AI adoption as workflow, governance, and change management.
- Can talk about regulated buyers without improvising security claims.
- Sees opportunity creation as account strategy, not activity volume.

Why This Role

I am interested in this role because it is a Strategic BDR role, not a generic SDR seat. The job sits close to account strategy: interpreting product signal, building account theses, mapping stakeholders, doing high-quality discovery, and partnering with Account Directors on deal progression. That matches both sides of my background. I have the sales reps from Airwallex and Boxo, and I have the AI-native GTM systems work from Salescraft and GTM Workspace. I can contribute quickly, but I also know OpenAI's enterprise motion is unique enough that I want to learn it from the inside.

Why SDR

I see it as a deliberate step into a higher-caliber room, not a retreat. There is a title step-back in the narrow sense because I have carried sales responsibility before. But the scope of learning is a step forward: OpenAI's product surface, enterprise buyers, regulated adoption patterns, and Account Director partnership are all things I want to learn properly. I am comfortable doing the work because I have done the hard parts of sales before: research, outbound, discovery, handoffs, stakeholder mapping, and follow-up.

Short version:

I am optimizing for slope, not title. I have sold before, but I want to learn OpenAI's enterprise motion correctly, contribute where I can, and earn the next level through performance rather than assumption.

Pushback: Will You Be Bored?

No. The only version that would bore me is pure volume without judgment. This role, as I understand it, is about product signal, account strategy, stakeholder mapping, and enterprise readiness. I like the craft of turning a vague account signal into a useful opportunity for a senior seller.

Pushback: Are You Overqualified?

I do not think about it that way. I have some useful reps, especially in regulated GTM and stakeholder mapping, but I do not assume that makes me qualified to skip learning OpenAI's enterprise motion. The useful part of my background is that I can bring judgment and ownership from day one while still being coachable.

AI Interest Story

My interest in AI became serious when I saw the same GTM failure loops repeating manually. The bottlenecks were not only effort. They were weak account selection, shallow research, generic outreach, poor handoffs, and bad CRM hygiene. I started building source-backed AI workflows around those problems: market mapping, enrichment, scoring, account intelligence, outbound drafts, and seller handoffs, with human review before anything customer-facing went out.

That evolved into broader AI workflow work: official sources, searchable memory, current-state tracking, workflow scripts, automations, and review gates. The lesson I took from that is that AI becomes useful when it is attached to a workflow, trusted sources, and a human decision path. That is why OpenAI's enterprise GTM problem is interesting to me. The buyer question is not "can the model do something impressive?" It is "where does this fit into a real workflow safely enough to scale?"

If Asked: Did You Use AI To Prep?

Yes. I used AI heavily, but in a review-first way. I used it to index notes, extract every question, build a checklist, pressure-test answers, and identify gaps. I did not treat it as a source of truth. I still check source material, separate official facts from interview intelligence, and keep a clean version for screen sharing.

Product Fluency

What Does OpenAI Do For Enterprises?

OpenAI helps enterprises use frontier AI in two broad ways. First, employees and teams use ChatGPT surfaces for knowledge work, analysis, writing, research, and internal productivity under enterprise controls. Second, developers and product teams use the API, tools, agents, and Codex to build AI into products, operations, and software workflows. The sales work is turning scattered usage, pilots, and technical enthusiasm into governed deployments with a clear workflow, stakeholder map, risk posture, and business outcome.

Explain An LLM Simply

A large language model is software trained on a huge amount of text and other data so it can predict and generate useful responses. It does not know things like a person does, and it can be wrong, but it is very useful when you give it context, instructions, tools, and a way to check its work.

Explain An Agent Simply

An agent is an AI system that can use a model plus tools and context to work through a task, not just answer one message. In enterprise settings, the important questions are what tools it can access, what permissions it has, what it can do without approval, how it is evaluated, and where a human reviews the output.

What Is RAG?

RAG stands for retrieval-augmented generation. It means the AI retrieves relevant information from a knowledge source, like company documents or a vector database, and uses that retrieved context when generating an answer. Commercially, it helps with internal or current knowledge because the model is not relying only on what it learned during training. It improves grounding, but it does not guarantee perfect accuracy.

Can Hallucination Be 100 Percent Avoided?

No. I would not promise 100 percent avoidance. The right answer is mitigation: better grounding, retrieval, tool calls, constraints, structured outputs, evals, monitoring, and human review for high-stakes workflows.

Strong line:

For regulated buyers, I would not sell certainty. I would sell controlled uncertainty: scope the workflow, ground the answer, measure behavior, and keep humans in the loop where the consequence of error is high.

RAG Versus Fine-Tuning

RAG changes what the model can look up at answer time. Fine-tuning or model optimization changes how the model behaves on a repeated task. If the issue is "the model needs current internal knowledge," start with retrieval. If the issue is "we need consistent behavior, format, classification, or style on a repeated task," model optimization may be relevant, but availability and recommended approach should be checked against current OpenAI docs.

Copilot Objection

That is a good sign, not a reason to stop the conversation. It means the organization is already investing in AI adoption. I would not lead with rip-and-replace. I would ask which workflows Copilot has solved, which ones remain unsolved, and where the buyer needs something beyond a bundled productivity assistant.

Good discovery questions:

- Where is Copilot already working well?
- Which teams are still building their own AI workflows?
- Are developers experimenting separately with APIs or agents?
- What governance, measurement, or workflow gaps remain?

Anthropic Objection

If a prospect is already using Anthropic, I would respect that. I am not there to attack a competitor or force rip-and-replace. I would ask what use cases Anthropic is supporting, what is working, and where there are unsolved workflows. OpenAI may land in a different part of the stack: employee productivity, developer workflows, voice/multimodal experiences, Codex, agents, or a specific enterprise deployment where OpenAI's product surface is a better fit.

Sales Method

Consultative And Value-Based Selling

Lead with the customer's priorities, pain points, workflow, stakeholders, and business outcome. Do not lead with "our product is superior" or cheaper pricing.

Answer spine:

1. What is the business priority?
2. What workflow is painful or strategically important?
3. Who owns the outcome?
4. What stakeholders can block or approve it?
5. What value metric matters?
6. What proof or next step makes the opportunity real?

Prospecting Motion

I run prospecting as a signal-led account strategy loop: define the account universe, find current evidence of AI urgency, map the right stakeholders, write one problem-first message, then learn from every reply or non-reply instead of treating outreach as a volume exercise.

Longer version:

I would start by defining the account universe around priority verticals: financial services, life sciences, retail, digital platforms, and large regulated enterprises in APAC. Then I would score accounts by AI readiness, business pain, likely enterprise value, and messageable trigger evidence: AI hiring, public AI strategy, developer/product signals, governance pressure, customer-support volume, internal knowledge complexity, or an executive transformation agenda. For top accounts, I would go deep: stakeholder map, current AI posture, likely governance concerns, and a specific workflow hypothesis.

AI In Prospecting

I use AI for the parts of prospecting where scale and synthesis help: building account briefs, summarizing annual reports, extracting stakeholder maps, identifying trigger signals, comparing messaging angles, and preparing drafts. But I keep human review on anything customer-facing. The standard is source-backed, specific, and respectful. AI should help me spend more time on judgment, not make me noisier.

APAC Nuance

I would not treat APAC as one prospecting motion. Singapore and ANZ are strong places for mature enterprise conversations: regulated buyers, technology-forward teams, English-language enterprise selling, and clearer governance discussions. India has huge scale and developer density, but the commercial conversation may need to be more disciplined around cost, ROI, and implementation value. SEA is more heterogeneous: the right channel, local partner, relationship path, and workflow proof can vary by country. My default would be to start with a market-specific account hypothesis, not a generic APAC sequence.

First APAC Market

If I had to pick one APAC market to land in first, I would choose Singapore. The role is based there, and it gives OpenAI a dense cluster of regulated financial institutions, regional headquarters, digital infrastructure, public-sector AI ambition, and English-language enterprise buyers. It is also a strong reference market: if you can show a governed, high-impact deployment in Singapore, that story can travel into broader SEA, ANZ, and regional enterprise conversations.

Deal Cards

Deal 1: Boxo Enterprise Platform Sale

Use for: complex sale, stakeholder mapping, multi-threading, trust-building, technical/compliance objections.

One-line thesis:

At Boxo, I learned enterprise sales the hard way: selling an early technical platform to a serious fintech buyer by making the deal underwriteable across executive value, product fit, engineering feasibility, and legal/compliance risk.

60-second walkthrough:

At Boxo, the best enterprise deal example was a fintech platform sale. We were early and selling miniapp infrastructure to a serious buyer, so the deal was not just a commercial pitch. The CEO cared about strategic upside and faster service launches. The product leader cared about whether it would become a real product experience. Engineering cared about SDK/API feasibility, auth and payment flows, and support ownership. Legal and compliance cared about vendor risk, data flow, SLAs, and liability.

My role was to hold those threads together: map stakeholders, translate objections back to product, keep the executive case tied to business value, and lead pricing plus legal/commercial negotiation. The result was a first enterprise proof point and the lesson that enterprise sales is risk removal, not just persuasion.

What this proves:

- Multi-threading across business, product, engineering, legal, and compliance.
- Value-based selling in a technical platform context.
- Honesty under enterprise buyer scrutiny.
- Relevance to OpenAI: enterprise AI adoption also needs executive appetite, workflow fit, technical validation, governance, and risk removal.

Deal 2: Airwallex Singapore GTM

Use for: quota, metrics, fintech workflow selling, disciplined outbound, coaching.

One-line thesis:

Airwallex proved I could execute in a mature, target-driven fintech GTM environment: define ICPs, sell workflow pain, close logos, and coach a repeatable outbound cadence.

Metrics:

- 200%+ annual quota attainment.
- 22 new logos.
- Average annual deal size to confirm before quoting precisely.

- Coached junior SDRs to around 150% target attainment.

60-second walkthrough:

At Airwallex, I joined as part of the Singapore outbound build-out. The core workflow pain was cross-border money movement: payments, FX, supplier payments, cards, spend, and finance operations. I built ICPs, account lists, messaging tests, discovery motion, follow-up discipline, and clean AE handoffs. The result was 200%+ annual quota attainment and 22 new logos. The bigger lesson was that GTM performance is an operating system: ICP, research, qualification, follow-up, CRM hygiene, coaching, and pipeline review.

Deal 3: Salescraft AI GTM Systems

Use for: AI in prospecting, systems from scratch, playbook built, failure/learning.

One-line thesis:

Salescraft was where I turned manual GTM work into AI-assisted account-intelligence systems: research, enrichment, scoring, outbound, CRM hygiene, and seller handoffs.

Metrics:

- 8 tech startups.
- USD 1M+ client pipeline.
- USD 200K revenue in 12 months.

Lesson:

AI should scale seller judgment, not replace it. The useful system has source evidence, routing state, confidence labels, approved messaging lanes, and human review.

Career Narrative

Timeline

Chapter	What it proves
Law / Cambridge	Systems lens, evidence discipline, institutional thinking.
Boxo	Enterprise risk removal, early-market selling, platform partnerships.
Airwallex	Quota discipline, mature fintech GTM, outbound cadence, coaching.
Salescraft	AI-native account intelligence and GTM systems.
GTM Workspace / AI workflow builds	Source-backed, review-first AI systems and

	workflow design.
OpenAI	Apply operator-builder arc to enterprise AI GTM at the frontier.

Transition Answers

Move	Clean answer
Law to startups	Law gave me a systems and evidence lens, but I wanted to work closer to products and operating systems that change behavior.
Boxo to Airwallex	Boxo taught me ambiguity. Airwallex let me test the same commercial muscles in a mature, quota-driven fintech environment.
Airwallex to Salescraft	I had lived the manual GTM loop and saw AI could turn it into systems: research, enrichment, scoring, outbound, CRM hygiene, and handoffs.
Builder path to OpenAI SDR	This is not abandoning building. It is choosing the highest-quality enterprise AI learning room, close to product signal, account strategy, and frontier deployment.

What Playbook Did You Build?

The most relevant playbook is my AI-native account-intelligence and enterprise outbound playbook. The premise is that in strategic accounts, every low-context touch spends trust, so the job is to find the few accounts and people where timing, workflow pain, and proof are strongest.

The playbook has five parts: define the account universe, capture source-backed signals, score by fit and timing, map stakeholders, and route the next best action to the seller. AI helps with research, enrichment, summarization, draft angles, and CRM hygiene, but humans own the commercial judgment and final outreach.

What Systems Did You Build From Scratch?

There are two categories. On the sales side, I helped build outbound and account-intelligence systems: ICP definition, account lists, signal capture, enrichment, scoring, messaging tests, CRM hygiene, and seller handoffs.

On the AI systems side, I built workflows around official sources, current state, durable memory, automations, and human review. The common theme is that I do not think of systems as dashboards. I think of them as operating loops: source, state, action, review, and learning.

What Worked / What Did Not

Area	What worked	What did not	What changed
Enterprise selling	Mapping stakeholders and risk owners explicitly.	Treating objections as persuasion problems.	Treat objections as the actual work to solve.
Outbound	Source-backed, timing-led outreach.	Generic AI-generated personalization.	Fix the ICP, premise, and signal routing before writing copy.
AI GTM	Evidence labels and human review.	Trusting generic databases or black-box scores.	Use source URLs, QA, confidence labels, and seller review.
AI products	Human review and schemas.	Letting the model own too much ambiguity.	Map workflow, edge cases, and trust boundaries before automating.

What Do You Still Need To Learn?

OpenAI's internal enterprise motion: how the best ADs define high-conviction opportunities, where deals stall, how product signal turns into durable deployment, and how technical, governance, and executive stakeholders are sequenced in the best accounts.

Behavioral Stories

Rapport / Trust

Use the Boxo enterprise fintech sale.

The rapport story is really a trust story. We were a young vendor selling to a serious fintech buyer. I had to show the executive sponsor that the business case was real, show the product leader that the experience could work, and make engineering feel their concerns were being handled rather than dismissed.

The trust-building move was to get specific. I worked through objections, brought technical concerns back to our product team, and avoided pretending we were more mature than we were. That created enough trust for the buyer to underwrite the risk.

Failure / Miss

Use Salescraft campaign learning.

Early on, it was tempting to treat AI prospecting as a copywriting problem: better personalization, more messages, more automation. What I learned is that if the list, signal, offer, or timing is wrong, AI just makes the wrong motion faster.

What changed afterward was my operating model. I started separating signal from scoring, requiring source evidence and confidence labels, routing weak signals to nurture instead of outbound, and diagnosing whether a miss came from list, offer, copy, deliverability, or timing before changing the message.

Ambiguity

Use Boxo or early AI workflow builds.

My default in ambiguity is to define the smallest useful system: what outcome matters, who owns it, what evidence exists, what is unknown, what next step creates learning, and what risk needs to be removed before we scale.

Speed And Quality

I move fast on reversible work and slow down on external, legal, pricing, compliance, or customer-impacting work. In AI workflows, that means drafts and research can move quickly, but customer-facing messages, claims, pricing, and sensitive actions need review.

Loom / Case Study Prep

Use this structure if assigned a strategic-customer project.

7-Minute Structure

Segment	Time
Why me / why OpenAI	1:00-1:30
Account point of view	1:00
One high-impact use case	1:30
How it could work technically	1:00
30-day engagement plan	1:00
Outbound message and standout close	1:00

Case Answer Spine

1. Account strategy: why this company, why now, what business priority matters.
2. One use case: specific workflow, target user, pain, measurable value.
3. Technical plausibility: channels, agent/workflow layer, tools/actions, customer systems, OpenAI surface, permissions, audit trail, human escalation.
4. Commercial plan: first champion, economic buyer, technical/governance validator, executive path.
5. Outbound: short, specific, tied to account evidence.
6. Risk: integration, data/security/privacy, adoption, accountability, incumbent vendors.
7. Why OpenAI: product breadth, quality, developer surface, enterprise controls, evals, and ability to scale from workflow wedge.

30-Day Engagement Plan

Week	Goal
1	Validate account hypothesis with first champion and current AI/workflow priorities.
2	Map economic buyer, technical/security validator, governance owner, and business owner.
3	Run workflow discovery around current process, data sources, risks, success metric, and pilot scope.
4	Package a sharp AD handoff: business problem, stakeholder map, pilot use case, risks, proof gaps, next meeting.

Outbound Message Template

Subject: AI workflow idea for [business priority]

Hi [Name], I noticed [current company signal]. It made me wonder whether your team is already looking at how AI could improve [specific workflow] without creating a governance headache.

One practical starting point could be [specific use case], where the goal is [business metric] while keeping [risk/control] explicit.

Worth a short conversation on whether this is a live priority for [team/company]?

Best,

Seth

Questions To Ask

Recruiter / HR

- What does great performance in this role look like in the first 90 days?
- Which parts of the role are most important early: account strategy, discovery quality, outbound quality, product fluency, or AD partnership?
- Are there any specific areas you would recommend I prepare more deeply for the next round?

SDR Manager

- What separates a strong Strategic BDR here from someone who is just good at activity volume?
- Where do the best BDRs add the most value to Account Directors?
- Which customer signals usually turn into the highest-quality opportunities?

AE / Commercial

- What makes a BDR handoff genuinely useful to you?
- Where do early OpenAI enterprise opportunities most often stall?
- What customer questions separate casual AI curiosity from real buying intent?

Product / Technical

- What level of product fluency do you expect from the best BDRs?
- Where do non-technical sellers usually overstate or misunderstand OpenAI's product surface?
- What customer governance questions should a BDR qualify before bringing in deeper technical resources?

Final Round

- What behaviors separate future OpenAI enterprise sellers from people who are just good at activity volume?
- What would you want this cohort to standardize in the first six months?
- Where have strong candidates underestimated the role?

Final Rehearsal Checklist

- [] 15-second intro is memorized.
- [] 30-second intro is natural.
- [] 2-minute intro can be cut down.
- [] Why OpenAI does not sound like brand worship.

- ☐ Why SDR sounds deliberate, humble, and high-conviction.
- ☐ RAG answer is clear in 30 seconds.
- ☐ Hallucination answer says "no, mitigate" confidently.
- ☐ Copilot and Anthropic answers are respectful and land-and-expand.
- ☐ Boxo deal card is under 3 minutes.
- ☐ Airwallex metrics are clean and exact enough.
- ☐ Salescraft AI GTM proof is concrete.
- ☐ Prospecting motion is signal-led, not volume-led.
- ☐ APAC answer avoids treating the region as one market.
- ☐ Case-study skeleton can be delivered in 7 minutes.
- ☐ No private notes, raw transcripts, recruiting-source notes, or internal repositories are visible.

Claims To Verify Before Live Use

- Current OpenAI product names, packaging, model names, security/privacy claims, and enterprise controls.
- Exact account facts for any named target company.
- Exact Airwallex average deal size, monthly close cadence, run-rate pipeline, close rate, and cycle length if asked.
- Which client names and metrics are comfortable to say aloud.